

Today we're going to study
Load balancing & routing games.

These are a good starting point because they
are fairly simple & they were one of the first
forays of CS into economics.

We'll warmup with load balancing games
before moving to the main event - routing games

Warmup: Load Balancing Games

These games are motivated by
data centers / server farms / super computers

... these are all situations where there are
a huge number of servers & a huge number of
jobs to run and somehow the system must
assign the jobs to the servers so that
performance is good.

The problem is that there are too many jobs/
servers to solve this centrally and so
each job "decides" its own assignment without
global information typically.

Hence, the jobs are playing a game.

The phrase "load balancing" comes from the fact
that you want to balance the assignment of
work (a.k.a. load) across the servers.

There's a simple model of this:

Here's a simple model of this:

we have m servers

& n jobs to run on the servers

→ job i has processing demand p_i
(size)

& can be served at
servers in S_i ← realistic!

An assignment A has $(i, j) \in A$ if

job i is run on
server j , & $j \in S_i$

Let $C(A)$ be the "cost" of assignment A

Goal: Find $A^* = \underset{A}{\operatorname{arg\min}} C(A)$

Depending on the setting

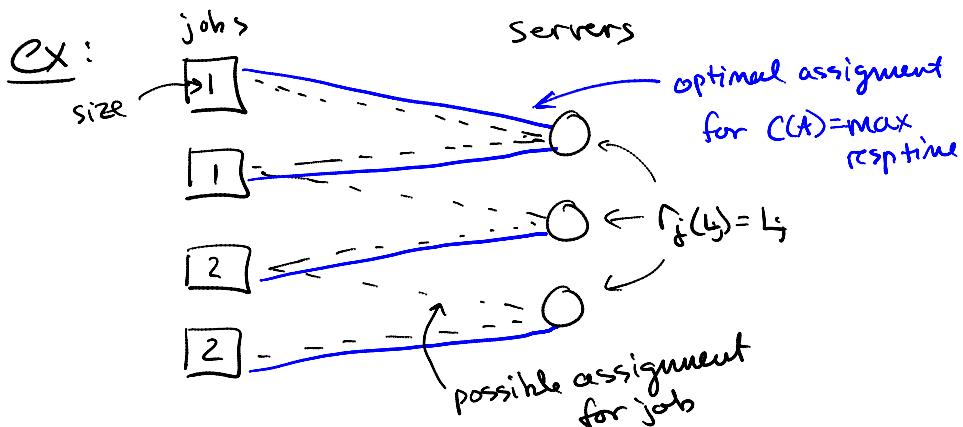
★ $C(A) = \max \text{"load"} = \max_j L_j$
where $L_j = \sum_{i:(i,j) \in A} p_i$

- $C(A) = \text{average load}$
 $= \frac{1}{m} \sum_j L_j$

★ $C(A) = \max \text{"response time"}$
 $= \max_j r_j(L_j)$ for inc $r_j(\cdot)$

- $C(A) = \arg \text{"response time"}$
 $= \frac{1}{m} \sum_j r_j(L_j)$

we'll do this



Centralized Approach

Centralized Approach

Given this model, what we want is

$$A^* = \underset{A}{\operatorname{arg\min}} = C(A)$$

→ this is NP-hard in general

Can't be found in cases where $n \& m$ are large



★ data centers
are large, often $n > 1,000,000$

Realistic Approach

"Greedy" → let each job greedily choose the server that minimizes its cost

i.e. job i chooses server j^* s.t.
 $j^* \in S_i \& j^* = \underset{j \in S_i}{\operatorname{arg\min}} r_j(L_j)$

★ So, the jobs are playing a game.

Players → jobs $S = \{1, \dots, n\}$
Actions → job i has actions $j \in S_i$
Payoffs → $r_j(L_j)$ for action $j \in S_i$

deterministic choice of a server
→ pure strategies

randomized decision
→ mixed strategies.

we'll focus on pure strategies
and leave randomized as a possible thw.

Note: This is a case where we're using game theory as a tool for analyzing an algorithm that doesn't initially sound like a "game".

Q: Is there a dominant strategy for the jobs?

A: No

Q: When is an assignment
a Nash Equilibrium?

A: If no job can improve cost by changing
its assignment

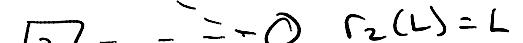
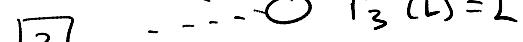
Specifically: A is a Nash equilibrium

$$\text{iff } \forall i, (i, j) \in A \Rightarrow \forall k \neq j, r_j(L_j) \leq r_k(L_k + p_i)$$

ex:  $r_1(L) = 2$

 $r_2(L) = 2$

Q: Which assignments are Nash equilibria?

A:   

ex:  $r_1(L) = L$

 $r_2(L) = L$

 $r_3(L) = L$

Q: Is the optimal assignment a Nash Equil.?

A: Yes!

Q: Are there other Nash Equil.?

A: No!

Okay, now we should all be
comfortable with the model.

Obviously it's a simplistic view of job assignment
in data centers, but it already has enough

features to make analysis interesting and
to provide some insight.

Q: What are some critiques of the model?

- A:
- jobs are assigned in 1 "batch"
 - jobs have same p_i at all servers
 - why would we arrive at a Nash?
- To learn about better models take my course next term - 147!

Q: Why is Nash Equilibrium a relevant concept?

A: Unclear... but maybe

jobs will "expect" others to do best thing
and so each job should "expect" others
to play the Nash Eg.

... putting aside the critiques though...

What do we want to understand about
this model?

- we'll focus on these
- 1) Does a Nash Equil. always exist?
 - 2) How efficient are the Nash Equil?
(how close to optimal is the resulting cost)
 - 3) How long does it take to find an equilibrium?

These questions will be our focus not just
for this model, but for most areas of
network economics.

1) Does there always exist an equilibrium?

→ Yes!

In particular, if we start with any assignment and make greedy moves (1 at a time), we'll arrive at an equilibrium

* This is called "Best Response Dynamics"

Q: What's the typical approach for showing that an algorithm like this will terminate?

(...and thus that we'll arrive at an equilibrium)

A: Find "something" that always decreases on each step, thus ensuring that there are no cycles. If there are no cycles then if there are a finite # of assignments the algorithm will eventually terminate

Q: What quantity always improves for a greedy move?

A1: $\max_j r_j(l_j) \rightarrow$ max response time.

... but this may remain the same, which causes problems.

A2: The sorted vector

$$v = (r_{c1}(l_{c1}), \dots, r_{cm}(l_{cm}))$$

$$\text{where } r_{c1}(l_{c1}) \geq \dots \geq r_{cm}(l_{cm})$$

we say $v_1 > v_2$ if

v_1 is lexicographically larger than v_2

$$\text{e.g. } (1, 2, 3, 4, 4, 5) > (1, 2, 2, 7, 8, 9)$$

T - .. ' - - - n

To see that a greedy move always decreases the sorted vector, note that if job i goes from $(j) \rightarrow (k)$

then 1) $(k) > (j)$

& 2) $r_{(j)}(L_{(j)}) > r_{(k)}(L_{(k)} + p_i)$

so $r_{(1)}(L_{(1)}) \dots r_{(j-1)}(L_{(j-1)})$ don't change

and $r_{(j)}(L_{(j)})$ will change to either

- current $r_{(j)}(L_{(j)} - p_i) \quad < r_{(j)}(L_{(j)})$
- current $r_{(j+1)}(L_{(j+1)}) \quad < r_{(j)}(L_{(j)})$
- current $r_{(k)}(L_{(k)} + p_i) \quad < r_{(j)}(L_{(j)})$

\Rightarrow the new vector is

lexicographically smaller.



So greedy steps will eventually terminate, and when they do, the assignment will be a Nash Equil.

Q: How long will it take to find a Nash Equil this way?

A: It might take exponential time...
... so we haven't made it any more efficient...

^

Note: In some cases finding an equil. can be done efficiently.

On to question 2:

Q: How "efficient" are Nash equilibria?

Before we can answer this we have to define how we'll measure "efficiency"...

how we'll measure "efficiency"...

$$\text{Let } A^* = \underset{A}{\operatorname{argmax}} C(A)$$

and let A^{ne} be a Nash equilibrium assignment.

We want to understand

$$\frac{C(A^{ne})}{C(A^*)}$$

but there are multiple Nash equil., so we have to be more careful.

There are 2 measures people use:

1) Price of Anarchy (PoA) of a game G

$$PoA(G) = \max_{A^{ne}} \frac{C(A^{ne})}{C(A^*)}$$

This PoA of a class of games \mathcal{G} is

$$PoA(\mathcal{G}) = \max_{G \in \mathcal{G}} \max_{A^{ne}} \frac{C(A^{ne})}{C(A^*)}$$

2) Price of Stability (PoS) of a game G .

$$PoS(G) = \min_{A^{ne}} \frac{C(A^{ne})}{C(A^*)}$$

and similarly

$$PoS(\mathcal{G}) = \max_{G \in \mathcal{G}} \min_{A^{ne}} \frac{C(A^{ne})}{C(A^*)}$$

Q: Any guesses about the PoS or PoA of our load balancing games?

Let's consider only $f_j(l_j) = L_j$ for now...

ex:



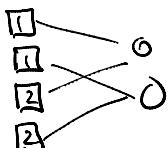
○

(every job can
go everywhere)



Q: What is the optimal assignment?

A:

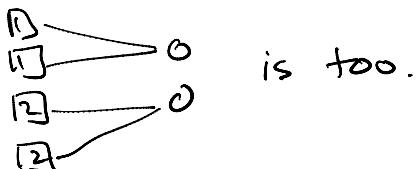


Q: Is this a Nash?

$$\overline{A}: \text{Yes!} \Rightarrow \text{PoS}=1$$

Q: Is this the only Nash?

A: No!



Q: What's the PoA?

A: The two Assignments we've seen are the only Nash equil.

So the 2nd is the worst

$$\Rightarrow P_0 A = \frac{4}{3}$$

Q: In this example the PoS was 1, does this hold more generally?

A: Yes! (for all LB games with $c(A) = \max_i r_i(A)$)

Consider A^* . No selfish move can increase $\max r_j(L_j)$, so

$\exists A^{ne}$ st $C(A^{ne}) \leq C(A^*)$, and
 $C(A^{ne}) < C(A^*)$ is impossible,
so $C(A^{ne}) = C(A^*)$.

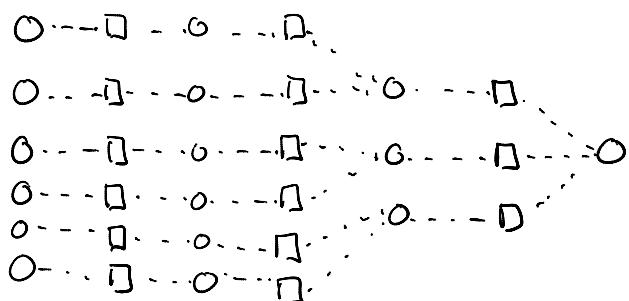
Q: What about the PoA? Can it be worse than $4/3$ in other load balancing games?

A: Yes!

One way \rightarrow makes $f_j(\cdot)$ steep
 e.g. set $f_j(l_{ij}) = L_j^{-\alpha}$
 for $\alpha \gg \infty$

Even with $r_j(l_j) = l_j$ we can make the Pot bad though:

Here's a way to get a PoA of 3:



All jobs are of size 1 & $L_j(L_j) = L_j$

2 Nash equil : assign left $\rightarrow C(R) = 1$
 assign right $\rightarrow C(A) = 3$

Clearly this can be generalized easily to get Pot of any integer value.

But: this took advantage of s_i not being equal.

Q: What if $s_i = s \forall i$? How large can the Pot be now?
(keep $g(L_i) = L_i$)

Thm: If $r_j(L_i) = L_i$ & $s_i = s \ \forall i$, then
 $PoA \leq 2 - \frac{2}{m+1}$

So, in this case, all Nash equilibria are nearly optimal.

nearly optimal.

It turns out that PoAs are surprisingly low in many cases... which is why game theory based algorithms have become popular.

Pf: Consider the worst Nash A^{ne}

Let server j^* have the highest load and job i^* have the smallest size among jobs at j^* .

WLOG : Assume ≥ 2 jobs at j^*
(else $PoA = 1$)

Claim 1: $P_{i^*} \leq \frac{1}{2} C(A^{ne})$

Claim 2: $\forall j \neq j^*, L_j \geq L_{j^*} - P_{i^*} \geq \frac{1}{2} C(A^{ne})$
(by def of Nash equil.)

Now, let's use these to bound $C(A^*)$

$$C(A^*) \geq \frac{\sum P_i}{M} = \frac{\sum L_i}{M} \geq \frac{C(A^{ne}) + \frac{1}{2} C(A^{ne})(m-1)}{M}$$

↑
Max \geq avg
↑
Sum of job sizes in the same ↑
above claims
↑
server j^* ↑
other $j \neq j^*$

$$\begin{aligned} \Rightarrow C(A^{ne}) &\leq \frac{m C(A^*)}{1 + \frac{1}{2}(m-1)} \\ &= \frac{2m}{m+1} C(A^*) \\ &= C(A^*) \left(2 - \frac{2}{m+1}\right) \end{aligned}$$

□

Q: Is this tight?

A: A possible HW question...

Okay, that's enough on load balancing games.

" we've seen that

- Nash equilibria always exist
- PoS = 1
- PoA can be large, but in some cases ($s_i = s$ & $f_j(L_j) = L_j$) it is very small.

Now on to the main event: Routing games

Topic 1: Routing Games aka Congestion Games

Now let's move on to the main event: Routing games.

Here the idea is that sources must choose their own routes through a network, and they do this to minimize their own latency.

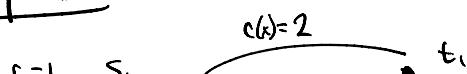
→ How much inefficiency does this create?

(compared to if they all decided together how to route)

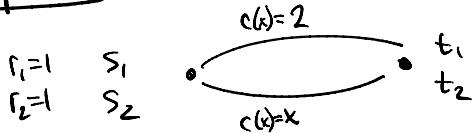
Non-CS example: Maybe an even better example of this is driving to work

- every day people choose a route to drive to work without knowledge of what others will pick, but how long it takes depends on the choices of others

Example 1:



Example +



Nash \rightarrow Both on bottom : Avg delay = 2
or 1 on top, 1 on bottom : Avg delay = $\frac{3}{2}$

OPT \rightarrow 1 on top 1 on Bottom : Avg delay = $\frac{3}{2}$

Defining the model

$$\text{Graph} = (V, E, C)$$

↑ each edge has a cost/latency function $c_e(x)$ that depends on the traffic x through the link. $c_e(\cdot)$ is non-neg, increasing, continuous.

Players : Each player has (s_i, t_i, r_i)
 ↑ ↑ ↑
 source destination amount of traffic

Actions : send r_i on a particular path from $s_i \rightarrow t_i$.
 (pure strategy again)

Payoff : Total latency of traffic:
 latency of sending x on Path p is
 $c_p(x) = \sum_{e \in P} c_e(f_e)$
 ↑
 flow on edge e
 (depends on other players)

Global Goal :

min avg latency

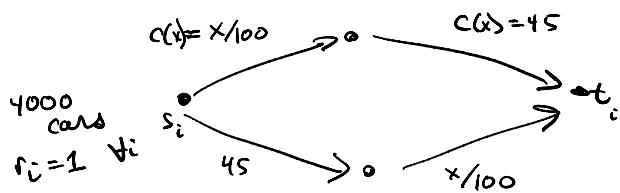
$$\Rightarrow \text{Cost} = \sum_p c_p(f_p) f_p$$

↑ ↑
 cost of path p flow on path p

$$= \sum_{e \in E} c_e(f_e) f_e$$

(like before we could choose many other goals here.)

Example 2 :



Q: What is the equilibrium?

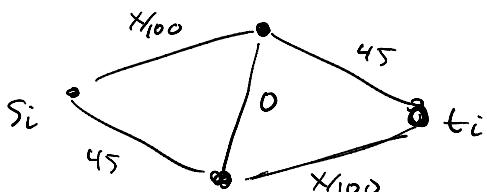
A: 2000 take top
2000 take the bottom

$$\text{Cost} = \frac{2000}{100} + 45 = 65$$

Q: Is this the optimal?

A: Yes.

Suppose we now add a freeway...



Q: What is the equilibrium now?

A:



$$\text{Cost} = \frac{4000}{100} + \frac{4000}{100} = 80$$

\Rightarrow The cost went up!

$$(PoS = \frac{80}{65} > 1!)$$

... can make PoS $\geq \frac{4}{3}$ with this example.

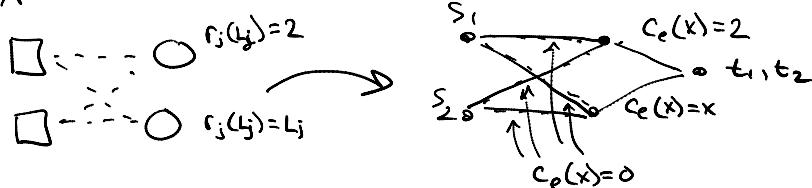
This is called Braess' Paradox

It is surprising that added capacity can hurt... but this happens frequently in real life (see ppt)

Example 3:

Q: How can we turn a load balancing game into a routing game.

A:



As with LB games, we'll discuss two things

- 1) Existence of Nash Equil
 - 2) Efficiency of Nash Equil
-

1)

We'll start by discussing the efficiency of the Nash Equilibrium for routing games...

We saw that in load balancing game, the PoA is unbounded in general, but is small (< 2) when $r_j(l_j) = l_j$ & $s_i = s$.

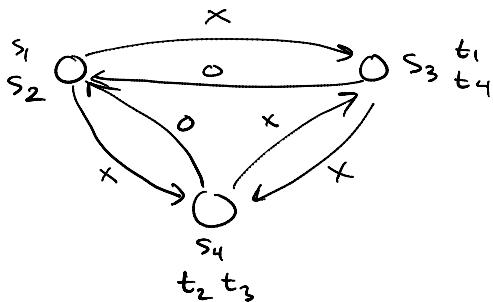
Again it is easy to see that its bad in general

in general C

But, what if $C_e(x) = ax + b$?

Q: Can you give me an example w/ $P(A) > 2$?

A:



Q: What is the opt way to route traffic?

A: the 1 hop paths cost = 4

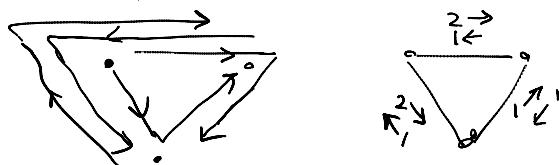
Q: Is this an equilibrium?

A: Yes

Q: Are there other equilibria?

A: Yes!

the 2 hop paths



cost = 10

$$P_0 A = \frac{10}{4}$$

★ This turns out to be almost the worst case!

Thm If all players have $r_i = r$ and $C_e(x) = ax + b$
 then $P_0 A \leq 2.5$

Then If $C_e(x) \neq ax+b$ then
 $P_{\text{of } A} \leq \phi^2 = \left(\frac{1+\sqrt{5}}{2}\right)^2 \approx 2.618$

golden ratio.

I can't resist making sure you all know about the golden ratio?

$$\phi = \frac{1}{1 + \frac{1}{1 + \frac{1}{1 + \dots}}} = \sqrt{1 + \sqrt{1 + \sqrt{1 + \dots}}}$$

$$= \frac{1 + \sqrt{5}}{2} = \lim_{n \rightarrow \infty} \frac{F(n+1)}{F(n)} \leftarrow \text{n-th Fibonacci} \quad \text{See ppt}$$

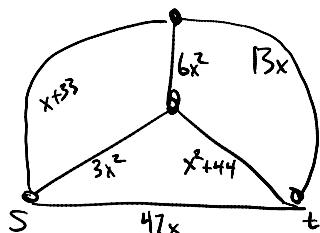
- ★ These are telling us (amazingly) that the network structure doesn't matter much for the efficiency of equilibria!

2) Existence of equilibria

Q: Do you think that there always exists an equilibria in these routing games?

A: No!

Here's an example w/ no equilibria



$r_1=1$ both have same $S \rightarrow t$
 $r_2=2$

To see that there is no equilibria:

Suppose 2 takes

then 1 takes

then 2 takes

then 1 takes

then 2 takes

T'll leave the details to you, first!

... I'll leave the details to you to check...

Q: So when does an equilibrium exist?

A: 1) if all flows have the same r_i

2) if all edges have $C_e(x) = \alpha_e x + b_e$

To prove this we will again use a variant/potential function.

* Let's start w/case of $r_i = 1 \quad \forall i$ (WLOG)

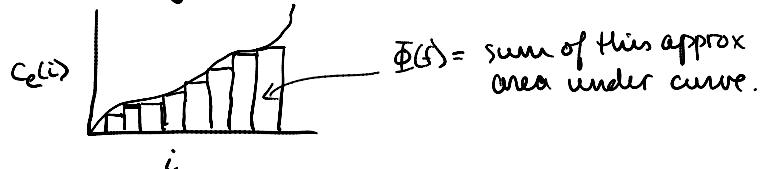
In particular, let's choose the variant f_{∞}

$$\Phi(f) = \sum_{e \in E} \sum_{i=1}^{f_e} c_e(i)$$

↑

total flow through graph

This may seem strange, but notice that



To see why this is a good variant for, let's consider the change to $\Phi(f)$ when one flow in f switches its route from $P \rightarrow P'$. Call the resulting flow f'

$$\begin{aligned}\Phi(f') - \Phi(f) &= \sum_{e \in P'/P} c_e(f_{e+1}) - \sum_{e \in P/P'} c_e(f_e) \\ &= C_P(f') - C_P(f)\end{aligned}$$

Change in Φ = change in flow payoff

Q: What does this property give us?

$\therefore f$ is even & odd \Rightarrow b. odd func

A: Greedy move + cost \Rightarrow \downarrow pot func

So, Any local minima of Φ is a Nash Equilibrium
since no player can make a move to $\downarrow \Phi$
 \Rightarrow no player can improve payoff
by deviating.

This is such an important property, it's given a name:

def A game is called a potential game if

\exists a potential function Φ st

$$\forall a, a': \Phi(a') - \Phi(a) = u_i(a') - u_i(a)$$

$u_i(x)$ is utility of player i (who changed more between a & a') for action profile x

In any potential game

- 1) \exists an equilibrium (the minima of Φ)
- 2) Best response dynamics will converge to an equilibrium.

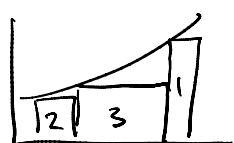
Okay, back to routing games...

the fact that they are a potential game (when $r_i=1$) means that an equilibrium exists & that Best response dynamics will converge.

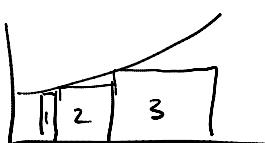
Q: What if r_i can be different?

Why doesn't Φ work anymore?

A: The order in the sum would now matter...



vs



... and the change in utility if i switches route always corresponds to having i last in the sum of Φ .

* Note that this is not a problem if

$C_e(a)$ is linear
 $\rightarrow \Phi$ is still a potential function in this case.

Another use for the potential function:

We know that the minimizer of Φ , a^g is a NE
... so if we can characterize the welfare of a^{ne} , we can bound the PoS.

Idea: If Φ is well "aligned" with the global cost function, then the PoS is small.

Thm: If $\forall a, \Phi(a) \leq C(a) \leq \alpha \Phi(a)$, then $PoS \leq \alpha$.

Pf: Let $a^{ne} = \operatorname{argmin}_a \Phi(a)$

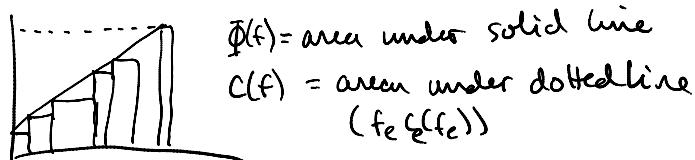
$$\begin{aligned} \text{Then } C(a^{ne}) &\leq \alpha \Phi(a^{ne}) \\ &\leq \alpha \Phi(a^*) \quad \leftarrow \text{optimizer of } C(\cdot) \\ &\leq \alpha C(a^*) \\ \Rightarrow PoS &\leq \alpha. \end{aligned}$$

□

Q: How can we apply this to routing games?

A: Consider $C(x) = ax + b$. ($b \geq 0$ since $C(x) \geq 0 \forall x$)

Then



$$C(f) \leq \Phi(f) \leq C(f)$$

$$\frac{1}{2}C(f) \leq \Phi(f) \leq C(f)$$

Okay, we've covered the basics of routing games

we've seen

- 1) When \exists an equilibrium
- 2) How bad equilibria can be

The conclusion is that in simple cases
equilibria will exist and be fairly
efficient, but in general equilibria
can be very inefficient.

So, we want to do something to make the
network more efficient.

(e.g. reduce traffic in our road network
example)

Q: What can we do?

A1: Add capacity... but Brayes Paradox
points out that this needs to be done
carefully.

A2: Add tolls, but how?

Consider the case of $\sum_i=1 H_i$.

Key idea of tolls \rightarrow

"cost" of a flow should be
not only latency, but also
Should include the added
cost to all other links caused

Cost to others being shared with.

* Principle of "marginal cost pricing"

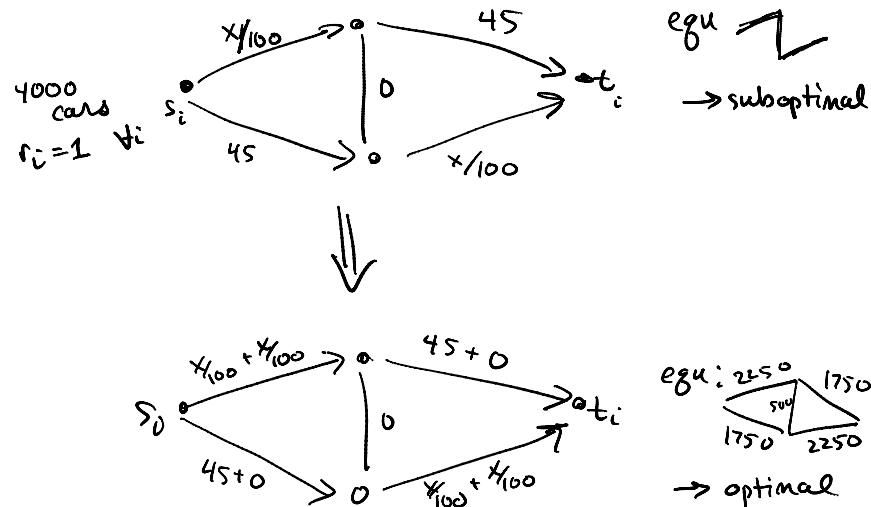
$$C_e(x) \rightarrow C_e^*(x) = C_e(x) + (x-1)(C_e(x) - C_e(x-1))$$

↑
 your latency
 ↓
 added latency
 to others being
 shared with.

Then: Marginal Cost Pricing guarantees that the PoS = 1.

Pf: On HW!

ex: Braess' Paradox.



That's it for routing games (for us).

As always, there's much more we could have covered ...

Important things to take away:

- game theory as algorithm design / analysis tool
- price of anarchy / price of stability
- potential functions / games
- conflict between local choices and global goal.
- principle of marginal cost pricing to align local choices with global goal.

Thanks!